

Coursera Data Science Capstone: Exploratory Data Analysis

Venkatesh Attinti

5/26/2020

Executive Summary

This milestone report is based on the exploratory data analysis of the swift key data provided in context of Data science capstone project. The data consist of 3 data file from different sources - (twitter,blogs,news). This report showcases the tidytext approach used for data analysis. More information regarding the tidy text approach can be accessed from here <https://www.tidytextmining.com/>

Data Summary

It is assumed that the data from <https://d396qusza40orc.cloudfront.net/dsscaphone/dataset/Coursera-SwiftKey.zip> is downloaded, unzipped and available in the working directory.

Below is the summary of the data loaded

```
library("stringi")

twitterRawData <- readLines("en_US.twitter.txt",warn=FALSE,encoding="UTF-8")
blogsRawData <- readLines("en_US.blogs.txt",warn=FALSE,encoding="UTF-8")
newsRawData <- readLines("en_US.news.txt",warn=FALSE,encoding="UTF-8")

datasummary <- data.frame("FileNames" =c("Twitter","Blogs","News"),
                          "FileSize"=c(format(object.size(twitterRawData), units = "MB", standard = "auto"),
                                         format(object.size(blogsRawData), units = "MB", standard = "auto"),
                                         format(object.size(newsRawData), units = "MB", standard = "auto")),
                          "FileLength"=c(length(twitterRawData),length(blogsRawData),length(newsRawData)),
                          "Wordcount"=c(sum(stri_stats_latex(twitterRawData)[4]),sum(stri_stats_latex(blogsRawData)[4]),sum(stri_stats_latex(newsRawData)[4])),
                          "NoOfChars"=c(sum(nchar(twitterRawData)),sum(nchar(blogsRawData)),sum(nchar(newsRawData))))

datasummary
```

##	FileNames	FileSize	FileLength	Wordcount	NoOfChars
## 1	Twitter	319 Mb	2360148	30451128	162096031
## 2	Blogs	255.4 Mb	899288	37570839	206824505
## 3	News	19.8 Mb	77259	2651432	15639408

Exploratory Data Analysis

In this section we will perform some exploratory data analysis using tidy data principles which is a powerful way to make handling data easier and more effective.

we will perform this analysis on the sample data set which is 2% of the original dataset.

Below are packages required to perform this analysis `library("tidyr") library("dplyr") library("tidytext") library("tm") library("openNLP") library("RWeka") library("tm")`

```
# Remove all non english characters as they cause issues down the road
twitterRawData <- iconv(twitterRawData, "latin1", "ASCII", sub="")
blogsRawData <- iconv(blogsRawData, "latin1", "ASCII", sub="")
newsRawData <- iconv(newsRawData, "latin1", "ASCII", sub="")

#sampling of the data set

twitterRawData_sample<- sample(twitterRawData,length(twitterRawData)*0.02)
blogsRawData_sample<- sample(blogsRawData,length(blogsRawData)*0.02)
newsRawData_sample<- sample(newsRawData,length(newsRawData)*0.02)

#write the sample files

dir.create("sampleDatafiles", showWarnings = FALSE)

write(twitterRawData_sample, "sampleDatafiles/twitterRawData_sample.txt")
write(blogsRawData_sample, "sampleDatafiles/blogsRawData_sample.txt")
write(newsRawData_sample, "sampleDatafiles/newsRawData_sample.txt")

remove(twitterRawData)
remove(blogsRawData)
remove(newsRawData)
```

Merging the sample data files into single corpus and then converting to tibble data frame format which will be used in all the further analysis

```
library("tidyr")
library("dplyr")
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library("tidytext")
library("tm")
```

```
## Loading required package: NLP
```

```
library("openNLP")
library("RWeka")
library("tm")
finalSampleData <- c(twitterRawData_sample,blogsRawData_sample,newsRawData_sample)
sampleData <- tibble(text = finalSampleData)
```

Pre-processing the data(invloves operations like removing the whitespaces, punctuation,stopwords, stemming etc) In tidy text the punctuations and converting to lower cases are automatically done during the unnesting tokens.

Unigrams

Unnesting tokens and removing the stopwords

```
data(stop_words)
tidySampleData <- sampleData %>% unnest_tokens(word, text) %>% anti_join(stop_words)
```

```
## Joining, by = "word"
```

```
#Removing whitespaces
tidySampleData$word <- gsub("\\s+", "",tidySampleData$word)
#Removing Numbers
tidySampleData<-tidySampleData[-grep("\\b\\d+\\b", tidySampleData$word),]

tidySampleData %>% count(word, sort = TRUE);
```

```
## # A tibble: 61,200 x 2
##   word      n
##   <chr> <int>
## 1 time    3458
## 2 love    3061
## 3 day     2910
## 4 people  2276
## 5 rt      1826
## 6 life    1558
## 7 lol     1468
## 8 happy   1314
## 9 night   1228
## 10 im     1209
## # ... with 61,190 more rows
```


using tidy tools,the word counts are stored in a tidy data frame.This allows us to pipe directly to the ggplot2 package

```
library(ggplot2)
```

```
##
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:NLP':
##
##   annotate
```

```
tidySampleData %>% count(word, sort = TRUE) %>% filter(n > 1000) %>% mutate(word = reorder(word, n)) %>%
```



Displaying the most common unnigrams using wordcloud

```
library(wordcloud)
```

```
## Loading required package: RColorBrewer
```


```
library("RColorBrewer")
dark2 <- brewer.pal(5, "Accent")
```

```
tidySampleData %>% count(word) %>% with(wordcloud(word, n, max.words = 100, rot.per=0.1, colors=dark2))
```



```
##      <chr>      <int>
## 1 happy birthday 184
## 2 mothers day   158
## 3 social media   95
## 4 happy mothers  93
## 5 ice cream      81
## 6 stay tuned     61
## 7 san diego      58
## 8 san francisco  58
## 9 friday night   53
## 10 beautiful day 50
## # ... with 172,542 more rows
```

```
bigrams_united %>% count(bigram, sort = TRUE) %>% filter(n > 35) %>% mutate(bigram = reorder(bigram, n))
```



Visualizing a Network of Bigrams with ggraph

It may be interesting to visualize all of the relationships among words simultaneously, rather than just the top few at a time. As one common visualization, we can arrange the words into a network, or “graph.”

```
library(igraph)
```

```
##
## Attaching package: 'igraph'
```

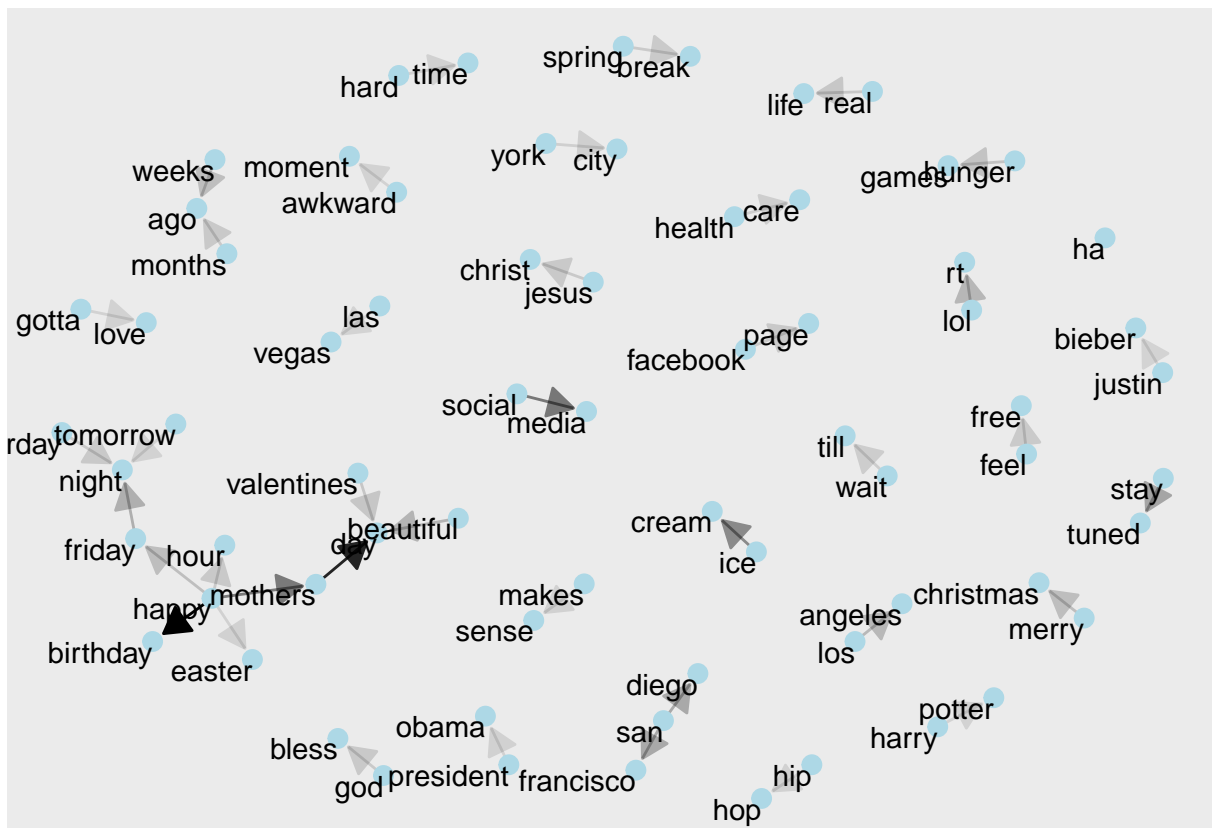
```
## The following objects are masked from 'package:dplyr':
##
##   as_data_frame, groups, union

## The following object is masked from 'package:tidyr':
##
##   crossing

## The following objects are masked from 'package:stats':
##
##   decompose, spectrum

## The following object is masked from 'package:base':
##
##   union
```

```
bigram_graph <- bigrams_filtered %>% count(word1, word2, sort=TRUE) %>% filter(n > 25) %>% graph_from_data_frame()
library(ggraph)
set.seed(123456)
a <- grid::arrow(type = "closed", length = unit(.15, "inches"))
ggraph(bigram_graph, layout = "fr") +
  geom_edge_link(aes(edge_alpha = n), show.legend = FALSE,
    arrow = a, end_cap = circle(.07, 'inches')) +
  geom_node_point(color = "lightblue", size=3) +
  geom_node_text(aes(label = name), vjust = 1, hjust = 1)
```



Trigrams – Tokenizing by 3-gram

```
library(tidyr)
stop_words <- rbind(stop_words, data.frame(word="amp", lexicon=""))
tidyBigramSampleData <- sampleData %>% unnest_tokens(bigram, text, token = "ngrams", n = 3)
#Seperating the bigram
tidyBigramSampleData_separated <- tidyBigramSampleData %>% separate(bigram, c("word1", "word2", "word3"))
bigrams_filtered <- tidyBigramSampleData_separated %>% filter(!word1 %in% stop_words$word) %>% filter(!word2 %in% stop_words$word) %>% filter(!word3 %in% stop_words$word)
#Removing whitespaces
bigrams_filtered$word1 <- gsub("\\s+", "", bigrams_filtered$word1)
bigrams_filtered$word2 <- gsub("\\s+", "", bigrams_filtered$word2)
bigrams_filtered$word3 <- gsub("\\s+", "", bigrams_filtered$word3)


bigrams_filtered$word1 <- gsub("\\\\'+", "", bigrams_filtered$word1)
bigrams_filtered$word2 <- gsub("\\\\'+", "", bigrams_filtered$word2)
bigrams_filtered$word3 <- gsub("\\\\'+", "", bigrams_filtered$word3)

#Removing Numbers
bigrams_filtered<-bigrams_filtered[-grep("\\b\\d+\\b", bigrams_filtered$word1),]
bigrams_filtered<-bigrams_filtered[-grep("\\b\\d+\\b", bigrams_filtered$word2),]
bigrams_filtered<-bigrams_filtered[-grep("\\b\\d+\\b", bigrams_filtered$word3),]

bigrams_united <- bigrams_filtered %>% unite(bigram, word1, word2, word3, sep = " ")
bigrams_united %>% count(bigram, sort=TRUE)
```

```
## # A tibble: 75,934 x 2
##   bigram                n
##   <chr>                <int>
## 1 happy mothers day    93
## 2 cinco de mayo        19
## 3 coffee coffee coffee  18
## 4 omg omg omg          18
## 5 st patricks day      16
## 6 ass ass ass           13
## 7 cake cake cake       12
## 8 greenville newspaper south 12
## 9 ha ha ha              12
## 10 happy valentines day  12
## # ... with 75,924 more rows
```

```
bigrams_united %>% count(bigram, sort = TRUE) %>% filter(n > 20) %>% mutate(bigram = reorder(bigram, n))
```

Plan of next steps

I have done the exploratory analysis. The next steps of this capstone project would be to finalize our predictive algorithm, and deploy our algorithm using shiny() app. As for the Shiny app it will consist of a simple user interface that will allow a user to enter text into a single textbox.